

JAX & Flax NNX Debugging Quick Reference

Core Challenge: The JIT Impact

- **JAX relies on** Compiles Python functions (tracing phase) into optimized code (execution phase).
- **Problem:** Standard Python `print()`, `pdb`, `breakpoint()` only see abstract *tracers* during tracing, not runtime *values* inside JIT-compiled code.
- **Goal:** Use JAX-specific tools or temporarily disable JIT to inspect runtime behavior.

Runtime Inspection (Inside JIT / Transformed Functions)

```
jax.debug.print("msg: {var}", var=value, ordered=True)
```

- **Purpose:** The JAX equivalent of `print()` for runtime values.
- **How:** Embeds print into the compiled graph.
- **Notes:** Use `ordered=True` for sequential output. Shows concrete values during execution.

```
jax.debug.breakpoint()
```

- **Purpose:** The JAX equivalent of `pdb.set_trace()` / `breakpoint()` for interactive debugging *inside* JIT.
- **How:** Pauses runtime execution, provides a `(jaxdb)` prompt.
- **Commands:** `p <var>` (print variable), `c` (continue), `q` (quit). Limited command set vs `pdb`.
- **Notes:** Can be made conditional using `jax.lax.cond()`. Standard Python `breakpoint()` can inspect *tracers* during tracing.

Enabling Standard Python Tools (Eager Execution)

```
jax.disable_jit()
```

- **Purpose:** Temporarily forces eager execution (like PyTorch/NumPy), making standard Python tools work.
- **How:**
 - Context Manager: `with jax.disable_jit(): ...` (Recommended)

- Global: `jax.config.update("jax_disable_jit", True)`
 - Env Var: `JAX_DISABLE_JIT=1`
- **Effect:** Standard `print()`, `pdb.set_trace()`, `breakpoint()`, IDE debuggers work directly on runtime values.
- **Drawback: Significant performance loss.** Use temporarily for debugging.
- **Limits:** May not fully step into `jax.vmap` / `jax.scan` internals.

Automatic Error Detection

`jax_debug_nans` Flag

- **Purpose:** Automatically find the source operation causing NaNs inside JIT.
- **How:** `jax.config.update("jax_debug_nans", True)` or `JAX_DEBUG_NANS=1`.
- **Effect:** If NaN detected, JAX re-runs eagerly to pinpoint the error source.
- **Limitations:** Slows execution; doesn't work with `jax.shard_map`; may flag intentional NaNs.

Flax NNX Specific Tools

`nnx.display(object)`

- **Purpose:** Visualize the structure and state of NNX objects (Modules, Optimizers, State).
- **How:** Provides a hierarchical view, uses `treescopes` for interactive display in notebooks/Colab if available.
- **Use:** Verify model architecture, check parameters/state initialization, inspect optimizer state. Analogous to `print(pytorch_model)`.

`Module.sow(variable_type, name, value, ...)`

- **Purpose:** Capture intermediate values (e.g., activations) during `nnx.Module.__call__` without altering function signatures.
- **How:** Stores `value` as an attribute (`name`) on the module instance. Access via `module.<name>.value`.
- **Notes:** Default appends to a tuple if called multiple times with the same name. Use `variable_type` (e.g., `nnx.Intermediate`) for organization.

Chex for Robustness & Debugging

Purpose: Library for reliable JAX code (assertions, testing).

Static Assertions: Check shapes, dtypes, ranks, structure (properties known at trace time).

- **Examples:** `chex.assert_shape`, `chex.assert_type`, `chex.assert_rank`, `chex.assert_trees_all_equal_shapes`.
- **Usage:** Place directly inside JIT-compiled functions.

Value Assertions: Check properties based on runtime values (NaN/Inf, closeness).

- **Examples:** `chex.assert_tree_all_finite`, `chex.assert_trees_all_close`.
- **Usage (in JIT):** Requires `@chex.chexify()` decorator.

@chex.chexify

- **Purpose:** Decorator to enable Chex *value* assertions inside JIT.
- **How:** Apply *outside* `@jax.jit`. Function returns `(error, result)`. Check `error.throw()` or `error.get()`.
- **Notes:** Performance overhead; debugging tool; **Doesn't work in Colab currently**. Usually need `chex.block_until_chexify_assertions_complete()` after call.

@chex.assert_max_traces (n=...)

- **Purpose:** Debug performance by detecting unexpected JIT re-compilations.
- **How:** Apply *inside* `@jax.jit`. Raises `AssertionError` if function is traced `> n` times.
- **Use:** Pinpoint causes of slow-downs due to unstable input shapes/types. Use `chex.clear_trace_counter()` if needed.

Monitoring & Visualization

TensorBoard

- **Purpose:** Visualize training metrics, images, profiling data.
- **Setup:** Create a Summary Writer (e.g., via `tensorflow`, `torch.utils.tensorboard`, or `tensorboardX`). Launch `tensorboard --logdir <log_dir>`.

- **Logging:** Use `writer.add_scalar`, `writer.add_image`, etc. **Remember `.item()` to convert JAX scalar arrays to Python scalars before logging.**
- **Profiling:** JAX profiler (`jax.profiler`) can output data viewable in TensorBoard.

Recommended Workflow Summary

1. **Static Checks:** `nnx.display`, Chex static assertions (`assert_shape`, etc.).
2. **Runtime JIT Issues:** Chex value assertions + `@chex.chexify` (for NaNs, etc.), `jax_debug_nans`, `jax.debug.print`, `jax.debug.breakpoint`.
3. **Complex Issues:** Temporarily use with `jax.disable_jit()`: to enable standard `pdb/print`.
4. **Performance:** `chex.assert_max_traces`, `jax.profiler` + TensorBoard.
5. **Monitor:** Use TensorBoard throughout.

Profiling

Profiling is also essential for understanding and improving your code. Xprof is a great tool for profiling JAX and Flax NNX, and is compatible with TensorBoard. [An excellent tutorial is available which includes profiling](#), so developers are encouraged to review that tutorial.

More Information

- JAX AI Stack - <https://jaxstack.ai>
- Chex - <https://chex.readthedocs.io>
- JAX - <https://jax.dev>
- Flax - <https://flax.readthedocs.io>